Lifelong Learning

CS 330

Plan for Today

The lifelong learning **problem statement**

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from **the meta-learning perspective**

A brief review of problem statements.

Multi-Task Learning

Learn to solve a set of tasks.

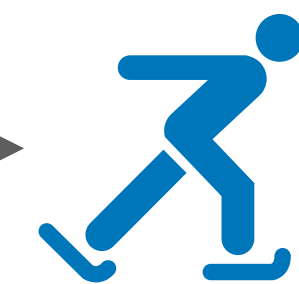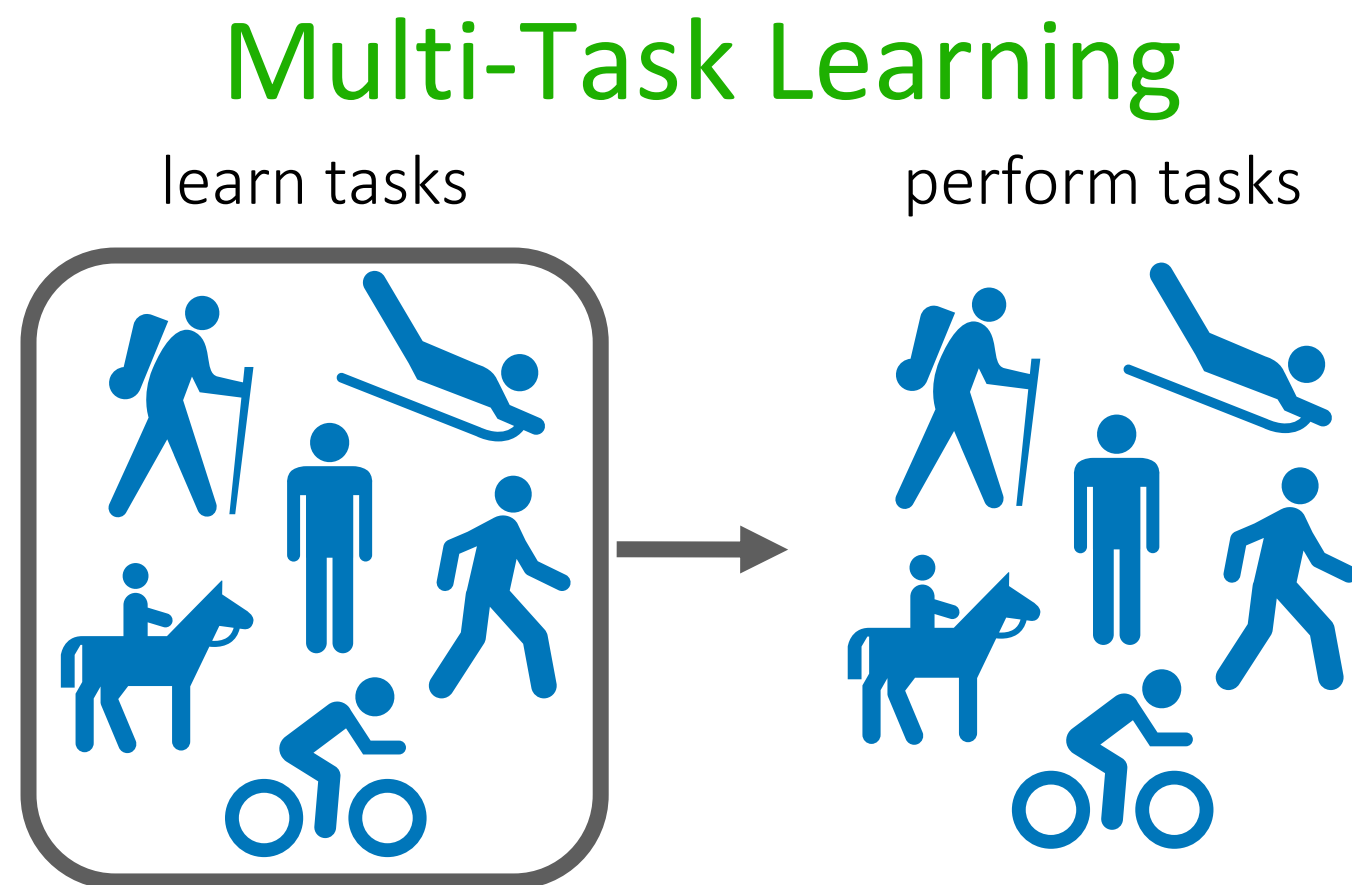learn tasks   perform tasks

Meta-Learning

Given i.i.d. task distribution, learn a new task efficiently
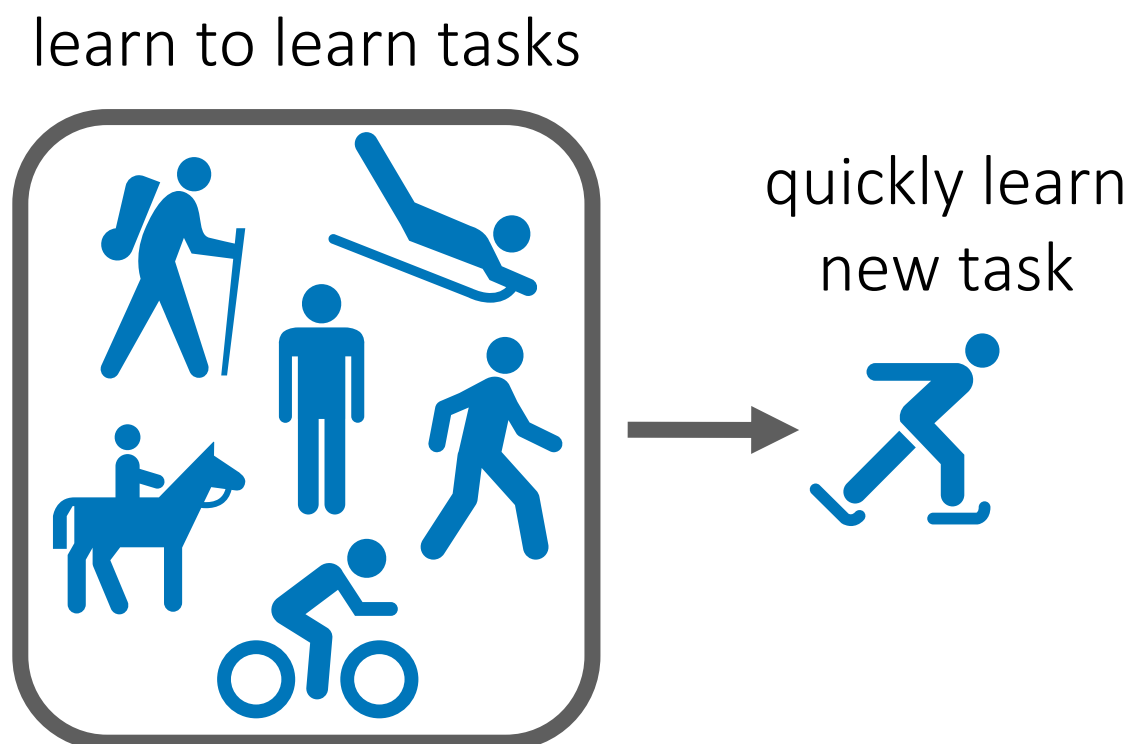
learn to learn tasks

quickly learn new task

Multi-Task Learning

learn tasks   perform tasks

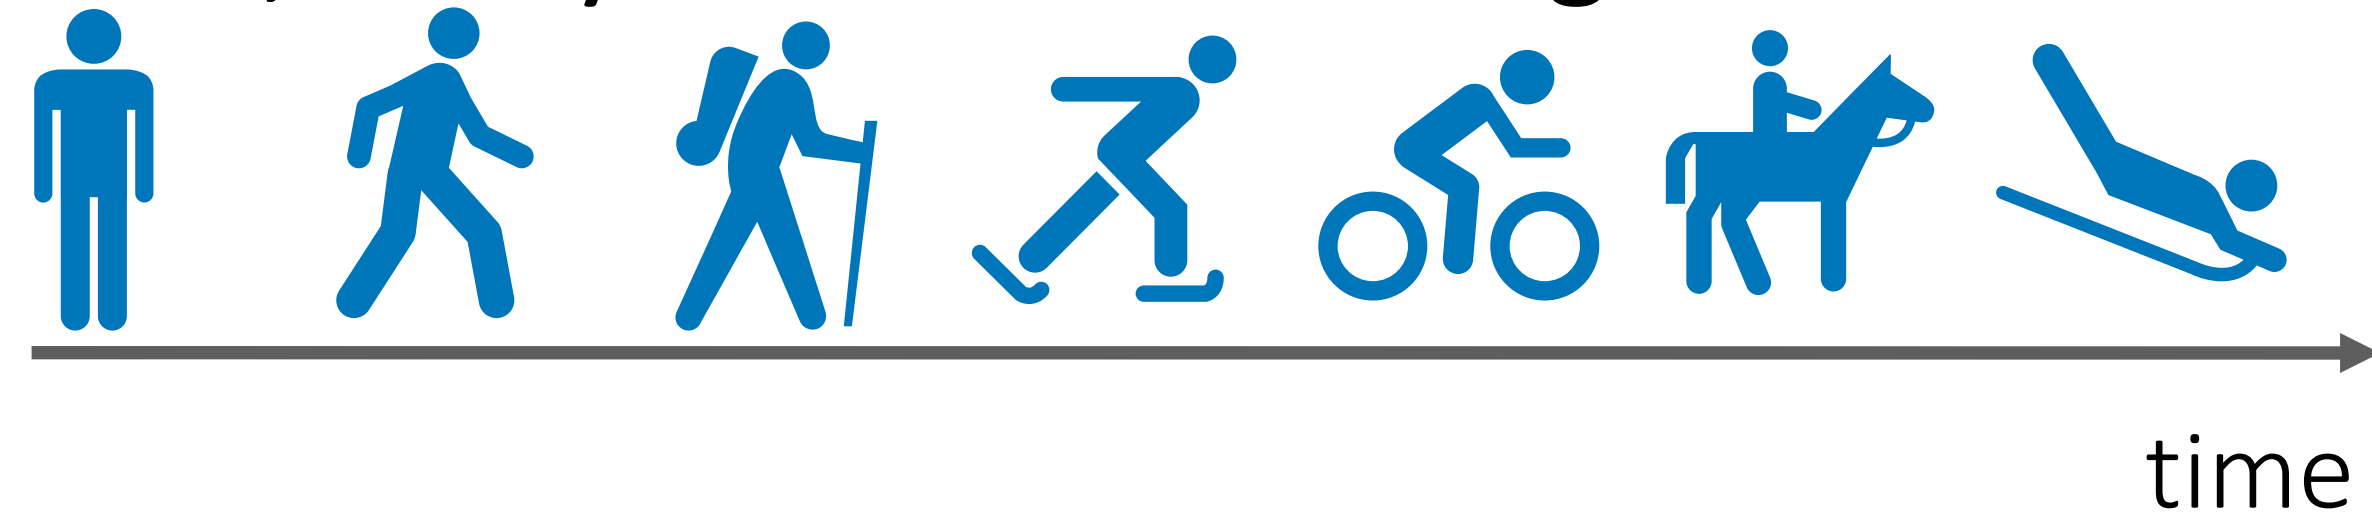## Meta-Learning

learn to learn tasks

quickly learn
new task

In contrast, many real world settings look like:

time

Our agents may not be given a large batch of
data/tasks right off the bat!

Some examples:
- a **student** learning concepts in school
- a deployed **image classification system** learning from a
  stream of images from users
- a **robot** acquiring an increasingly large set of skills in
  different environments
- a **virtual assistant** learning to help different users with
  different tasks at different points in time
- a **doctor's assistant** aiding in medical decision-making

4

Some Terminology

Sequential learning settings

online learning, lifelong learning, continual learning, incremental learning, streaming data

distinct from **sequence data** and **sequential decision-making**

What is the lifelong learning *problem statement*?

Exercise:

1. Pick an example setting.

2. Discuss problem statement in your break-out room:

   (a) how would you set-up an experiment to develop & test your algorithm?

   (b) what are desirable/required properties of the algorithm?

   (c) how do you evaluate such a system?

Example settings:

A. a **student** learning concepts in school

B. a deployed **image classification system** learning from a stream of images from users

C. a **robot** acquiring an increasingly large set of skills in different environments

D. a **virtual assistant** learning to help different users with different tasks at different points in time

E. a **doctor's assistant** aiding in medical decision-making

Desirable properties/considerations	Evaluation setup

What is the lifelong learning *problem statement*?

Problem variations:

 - **task/data order**: i.i.d. vs. predictable vs. curriculum vs. adversarial

 - **discrete** task boundaries vs. **continuous** shifts (vs. both)

 - **known** task boundaries/shifts vs. **unknown**

Some considerations:

 - model **performance**

 - data **efficiency**

 - **computational** resources

 - **memory**

 - others: privacy, interpretability, fairness, test time compute & memory

Substantial variety in problem statement!

What is the lifelong learning *problem statement*?

General [supervised] online learning problem:

for t = 1, ..., n

observe x_t <— if **observable task boundaries**: observe x_t, z_t

predict $\hat{y}_t$

observe label y_t

i.i.d. setting: $x_t \sim p(x), y_t \sim p(y|x)$ **streaming setting**: cannot store (x_t, y_t)

p not a function of t - lack of memory

- lack of computational resources

otherwise: $x_t \sim p_t(x), y_t \sim p_t(y|x)$ - privacy considerations

- want to study neural memory mechanisms

true in some cases, but not in many cases!

- recall: replay buffers

What do you want from your lifelong learning algorithm?

minimal regret (that grows slowly with t)

regret: cumulative loss of learner — cumulative loss of best learner in hindsight

$$Regret_T := \sum_1^T \mathcal{L}_t(\theta_t) - \min_\theta \sum_1^T \mathcal{L}_t(\theta)$$

(cannot be evaluated in practice, useful for analysis)

Regret that grows linearly in t is trivial.   Why?

What do you want from your lifelong learning algorithm?

positive & negative transfer

positive forward transfer: previous tasks cause you to do better on future tasks

compared to learning future tasks from scratch

positive backward transfer: current tasks cause you to do better on previous tasks

compared to learning past tasks from scratch

positive -> negative : better -> worse

Plan for Today

The lifelong learning **problem statement**

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from **the meta-learning perspective**

Approaches

Store all the data you've seen so far, and train on it.   —> **follow the leader** algorithm

+ will achieve very **strong performance**

- **computation intensive**   —> **Continuous fine-tuning** can help.

- can be **memory intensive**   [depends on the application]

Take a gradient step on the datapoint you observe.   —> **stochastic gradient descent**
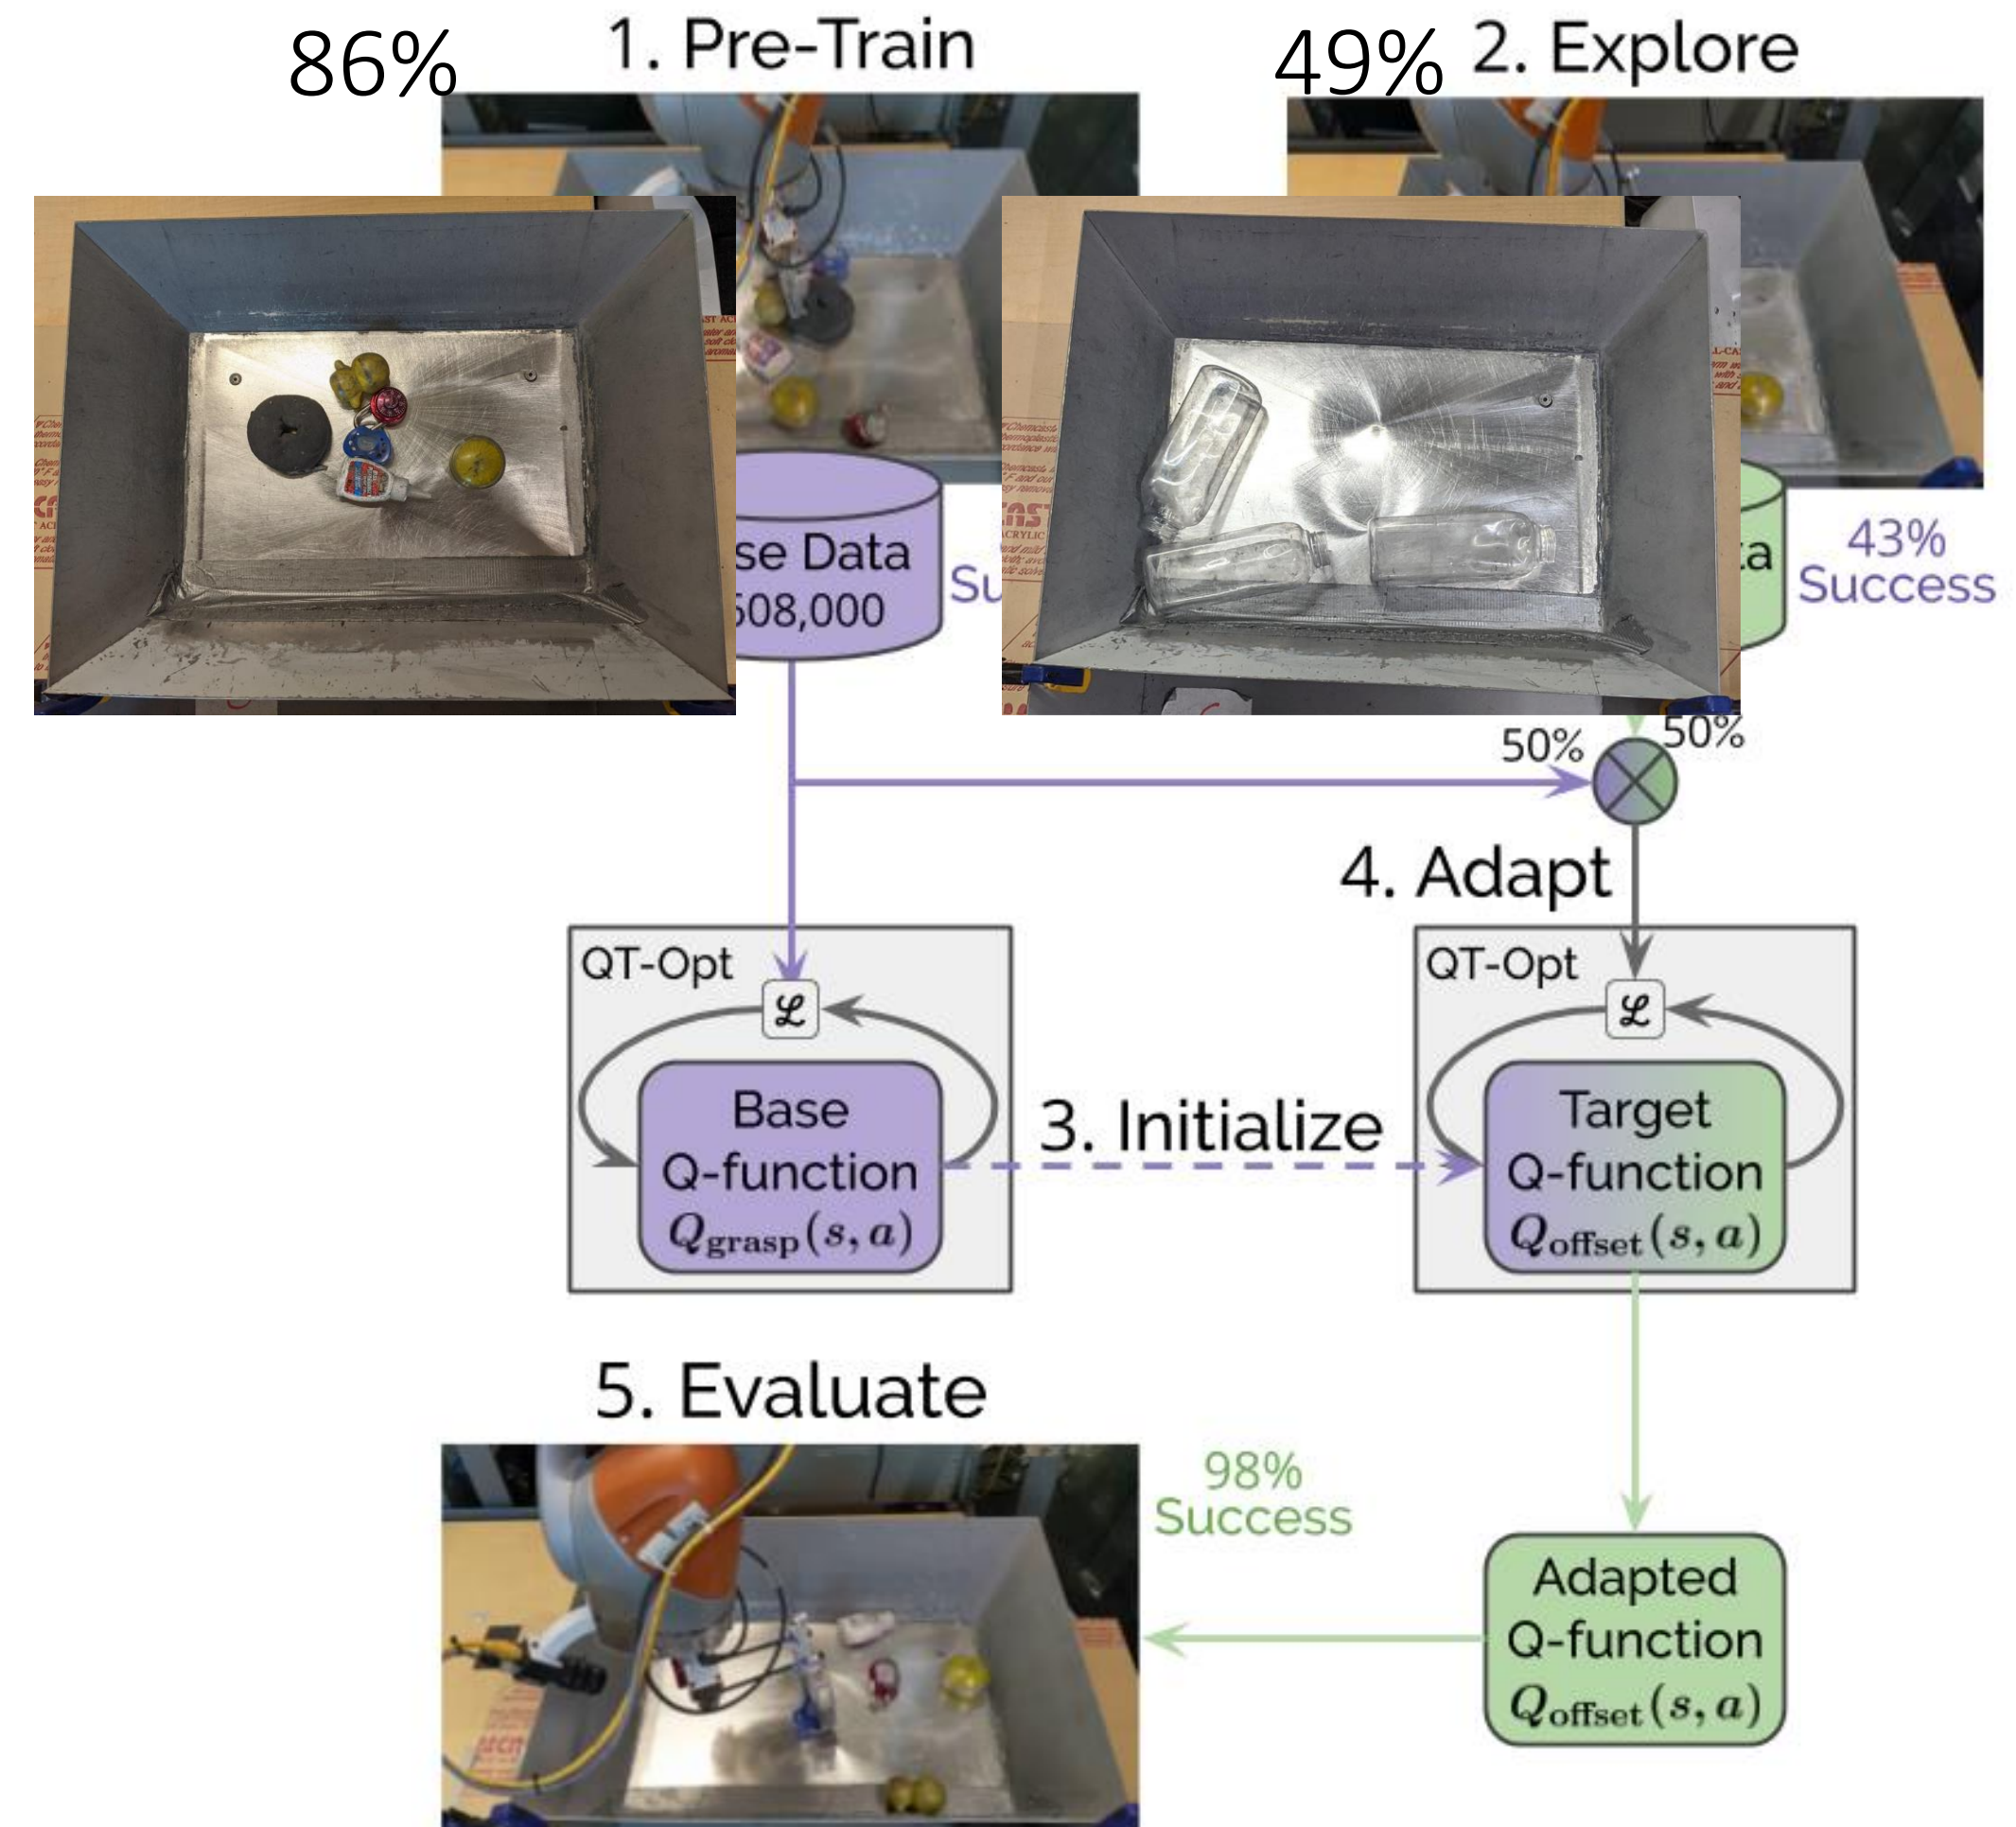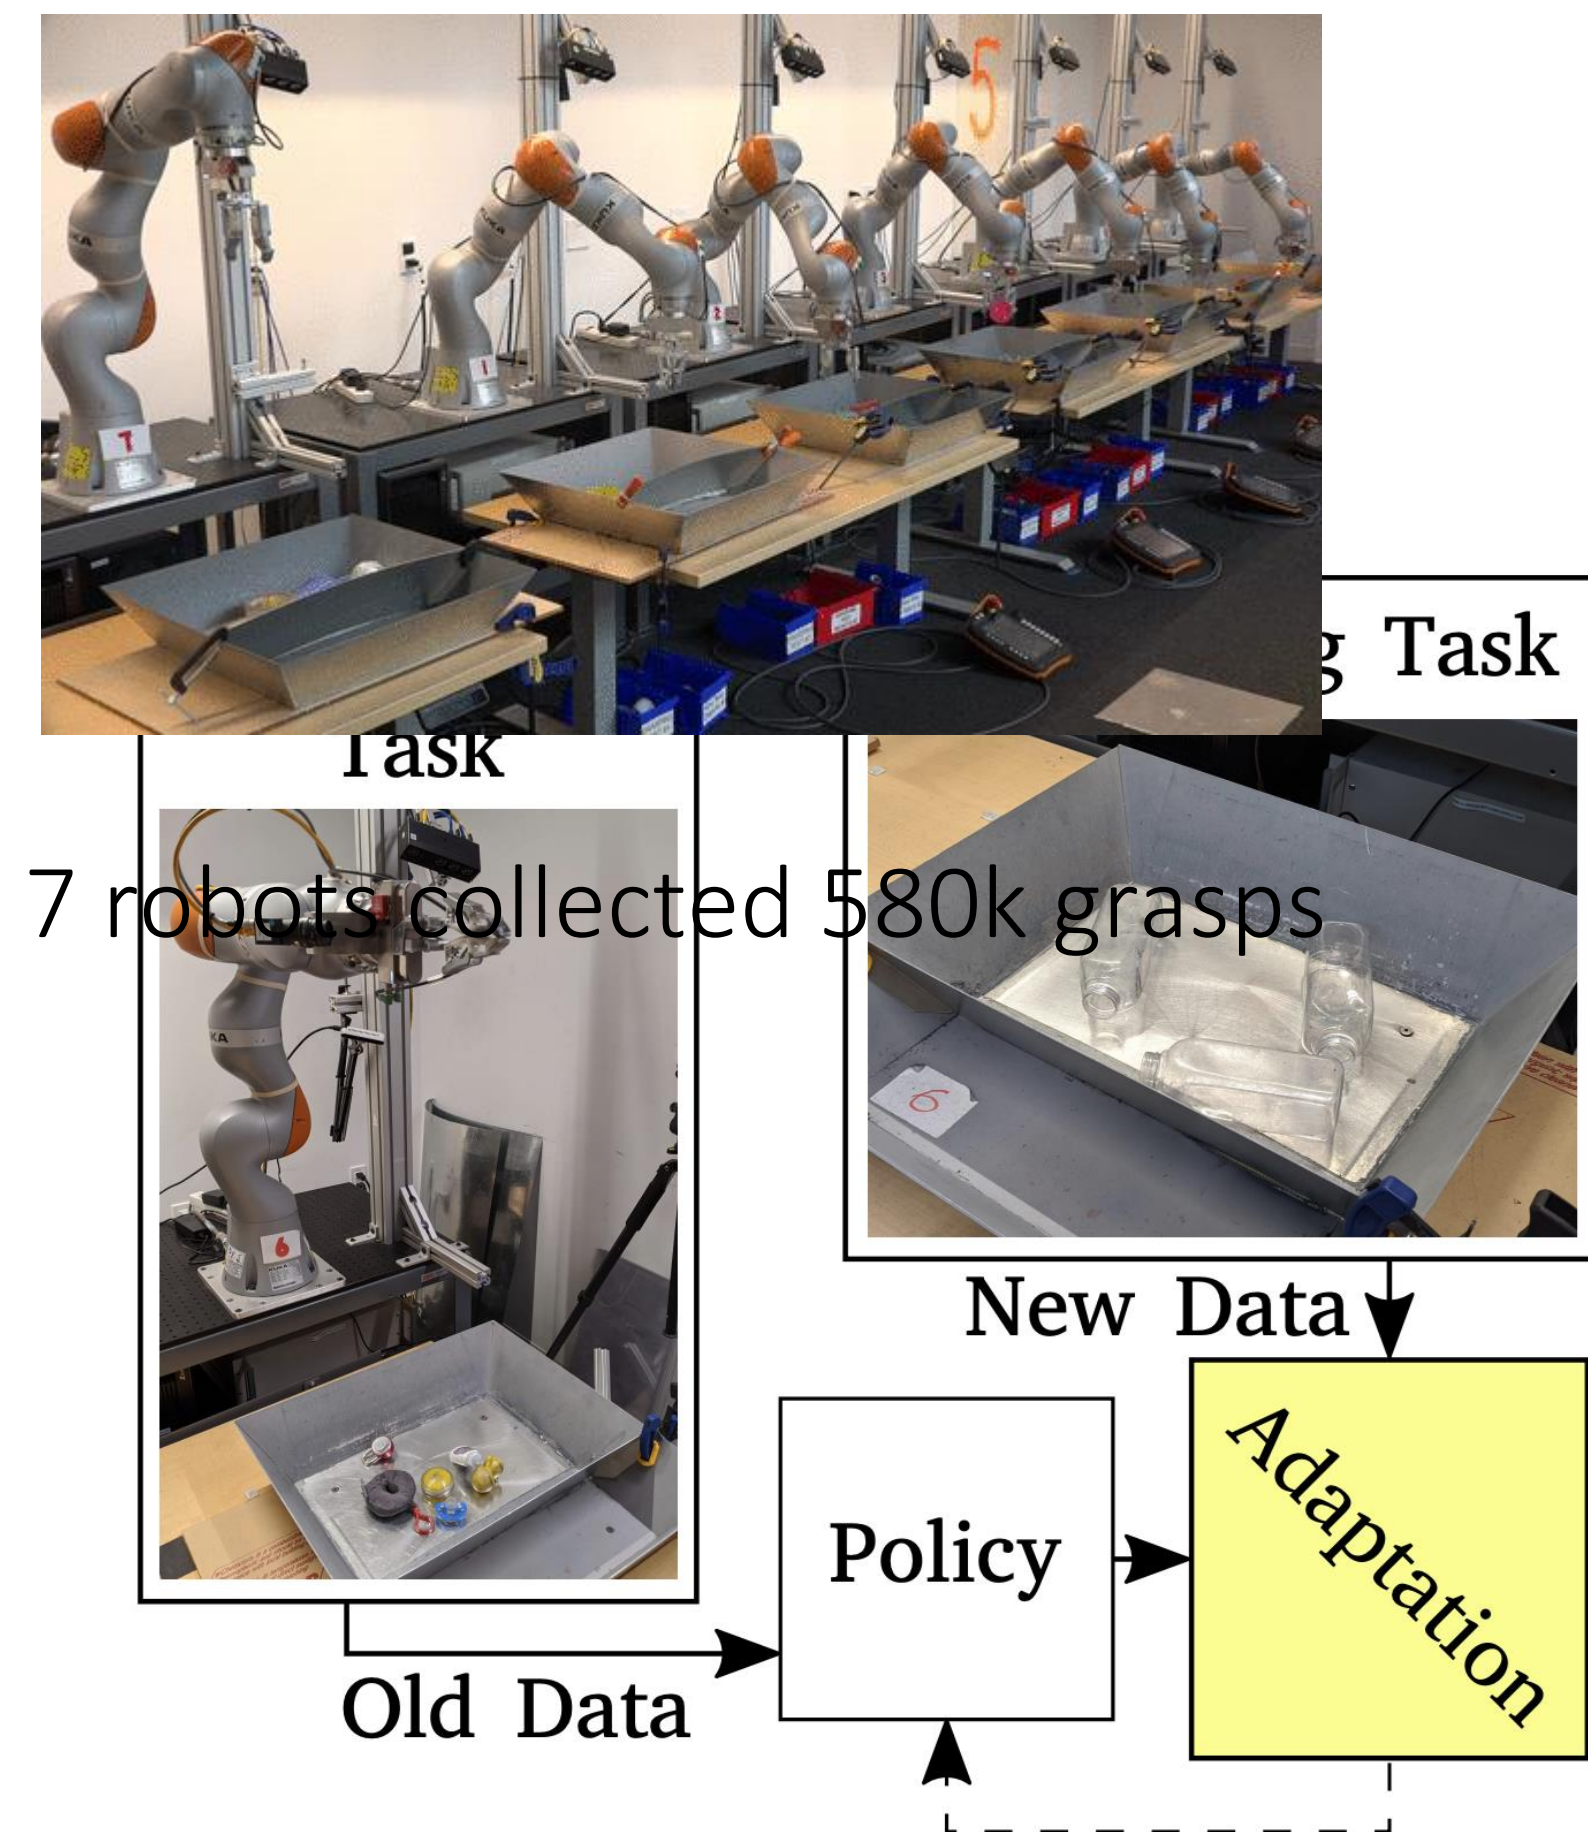
+ computationally cheap

+ requires 0 memory
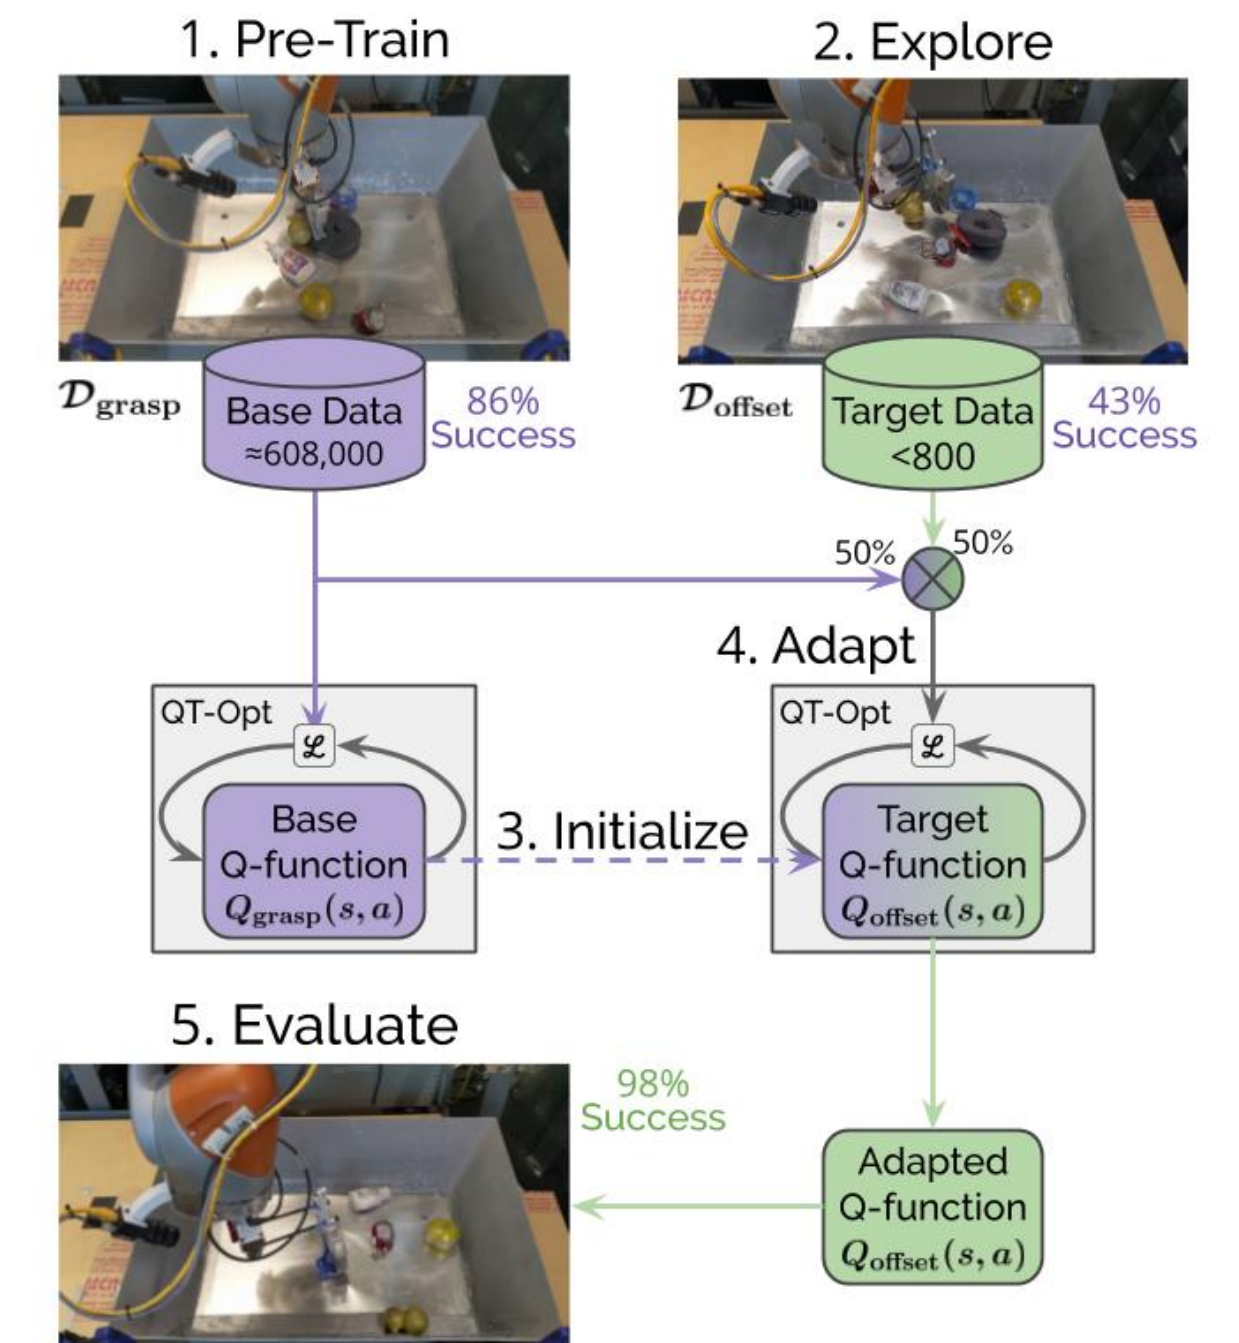
- subject to **negative backward transfer**   sometimes referred to as catastrophic forgetting
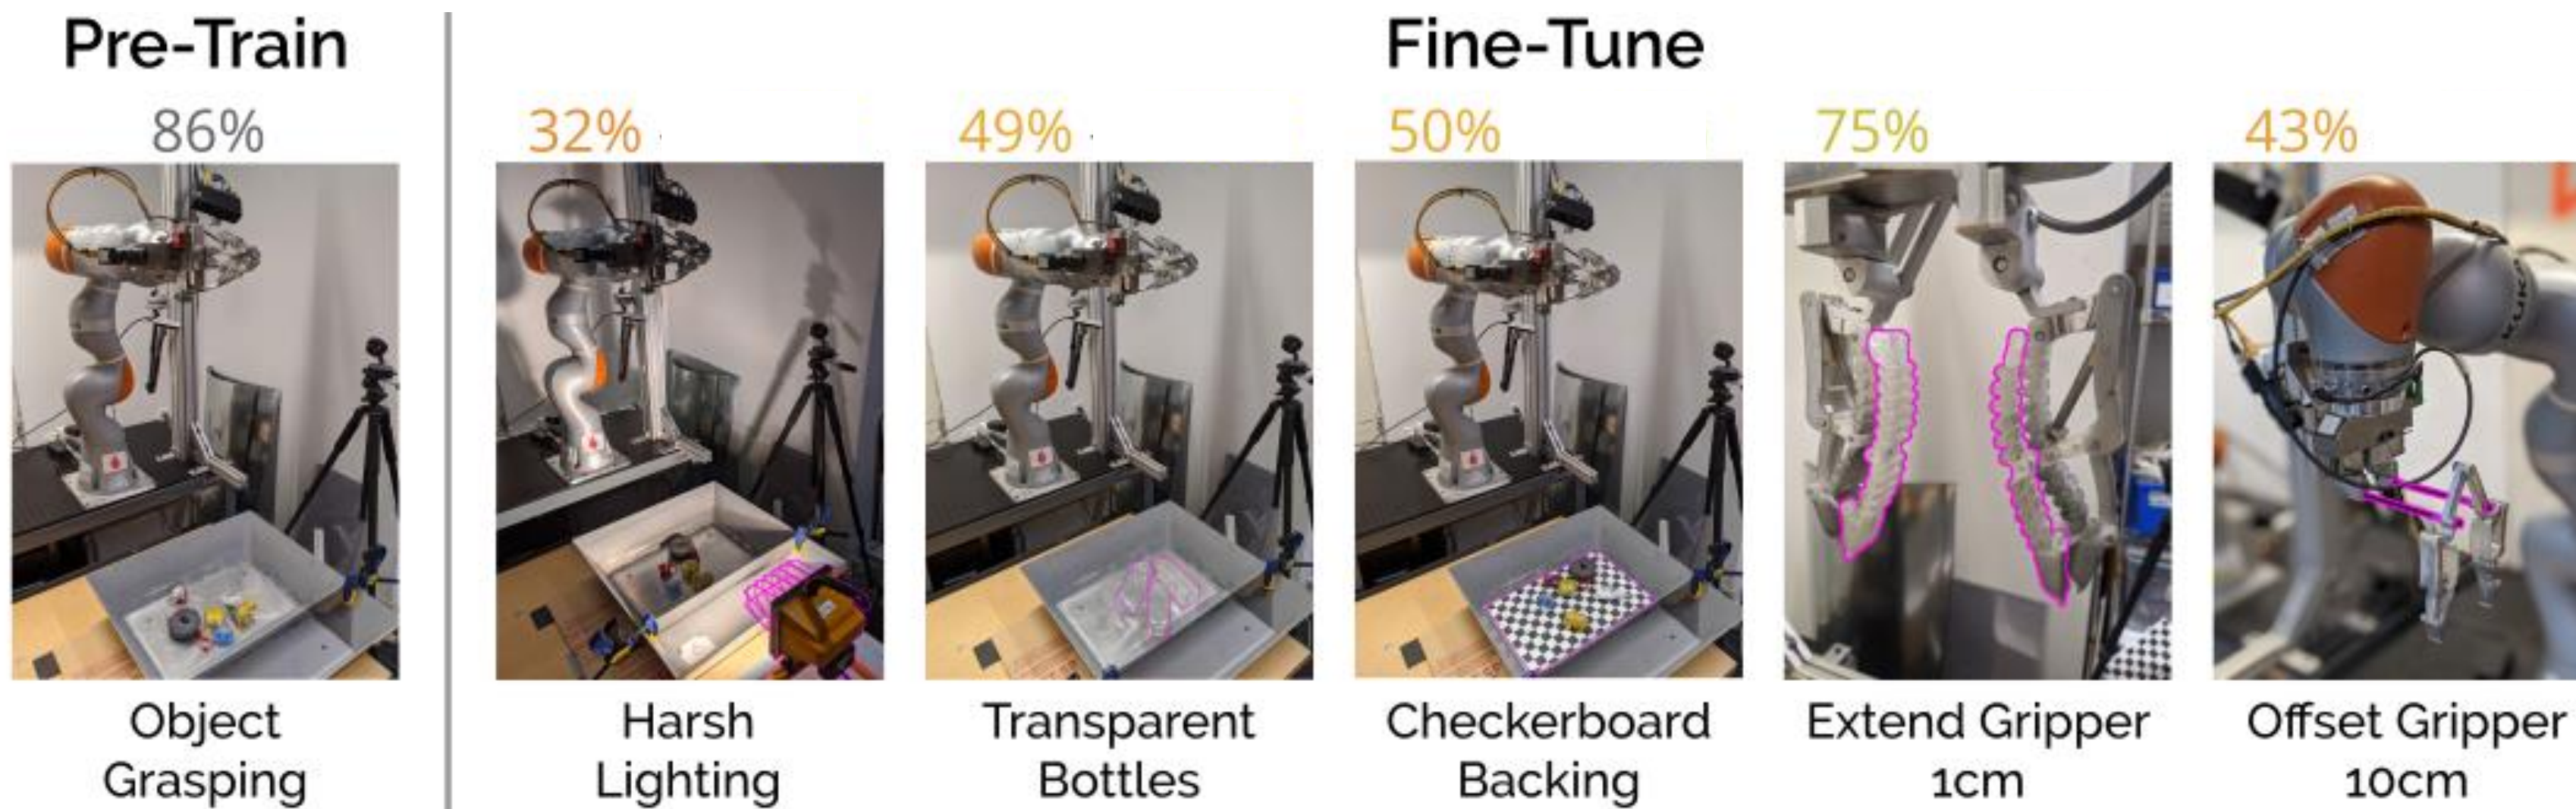
"forgetting"

- slow learning

Very simple continual RL algorithm

7 robots collected 580k grasps

1. Pre-Train — 86%
2. Explore — 49%
43% Success
50% | 50%
3. Initialize
4. Adapt
5. Evaluate — 98% Success

Base Q-function $Q_{grasp}(s,a)$
Target Q-function $Q_{offset}(s,a)$
Adapted Q-function $Q_{offset}(s,a)$

Base Data 608,000

Task / Task
Old Data → Policy → Adaptation
New Data

Julian, Swanson, Sukhatme, Levine, Finn, Hausman, Never Stop Learning, 2020

Very simple continual RL algorithm

Pre-Train
86%
Object Grasping

Fine-Tune
32% Harsh Lighting
49% Transparent Bottles
50% Checkerboard Backing
75% Extend Gripper 1cm
43% Offset Gripper 10cm

1. Pre-Train
$\mathcal{D}_{grasp}$ Base Data ≈608,000 — 86% Success

2. Explore
$\mathcal{D}_{offset}$ Target Data <800 — 43% Success

50% 50%

4. Adapt

QT-Opt — Base Q-function $Q_{grasp}(s,a)$

3. Initialize

QT-Opt — Target Q-function $Q_{offset}(s,a)$

5. Evaluate
98% Success

Adapted Q-function $Q_{offset}(s,a)$

Julian, Swanson, Sukhatme, Levine, Finn, Hausman, Never Stop Learning, 2020

Very simple continual RL algorithm

Can we do better?

Julian, Swanson, Sukhatme, Levine, Finn, Hausman, Never Stop Learning, 2020

Plan for Today

The lifelong learning **problem statement**

Basic approaches to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement
from **the meta-learning perspective**

Case Study: Can we modify vanilla SGD to avoid negative backward transfer?
(from scratch)

Idea:
(1) store small amount of data per task in **memory**

(2) when making updates for new tasks, ensure that they don't **unlearn** previous tasks

How do we accomplish (2)?

learning predictor $y_t = f_\theta(x_t, z_t)$ memory: $\mathcal{M}_k$ for task z_k

For $t = 0, \ldots, T$

minimize $\mathcal{L}(f_\theta(\cdot, z_t), (x_t, y_t))$

subject to $\mathcal{L}(f_\theta, \mathcal{M}_k) \leq \mathcal{L}(f_\theta^{t-1}, \mathcal{M}_k)$ for all $z_k < z_t$

(i.e. s.t. loss on previous tasks doesn't get worse)

Assume local linearity:

$$\langle g_t, g_k \rangle := \langle \frac{\partial \mathcal{L}(f_\theta, (x_t, y_t))}{\partial \theta}, \frac{\mathcal{L}(f_\theta, \mathcal{M}_k)}{\partial \theta} \rangle \geq 0 \qquad \text{for all } z_k < z_t$$

Can formulate & solve as a QP.

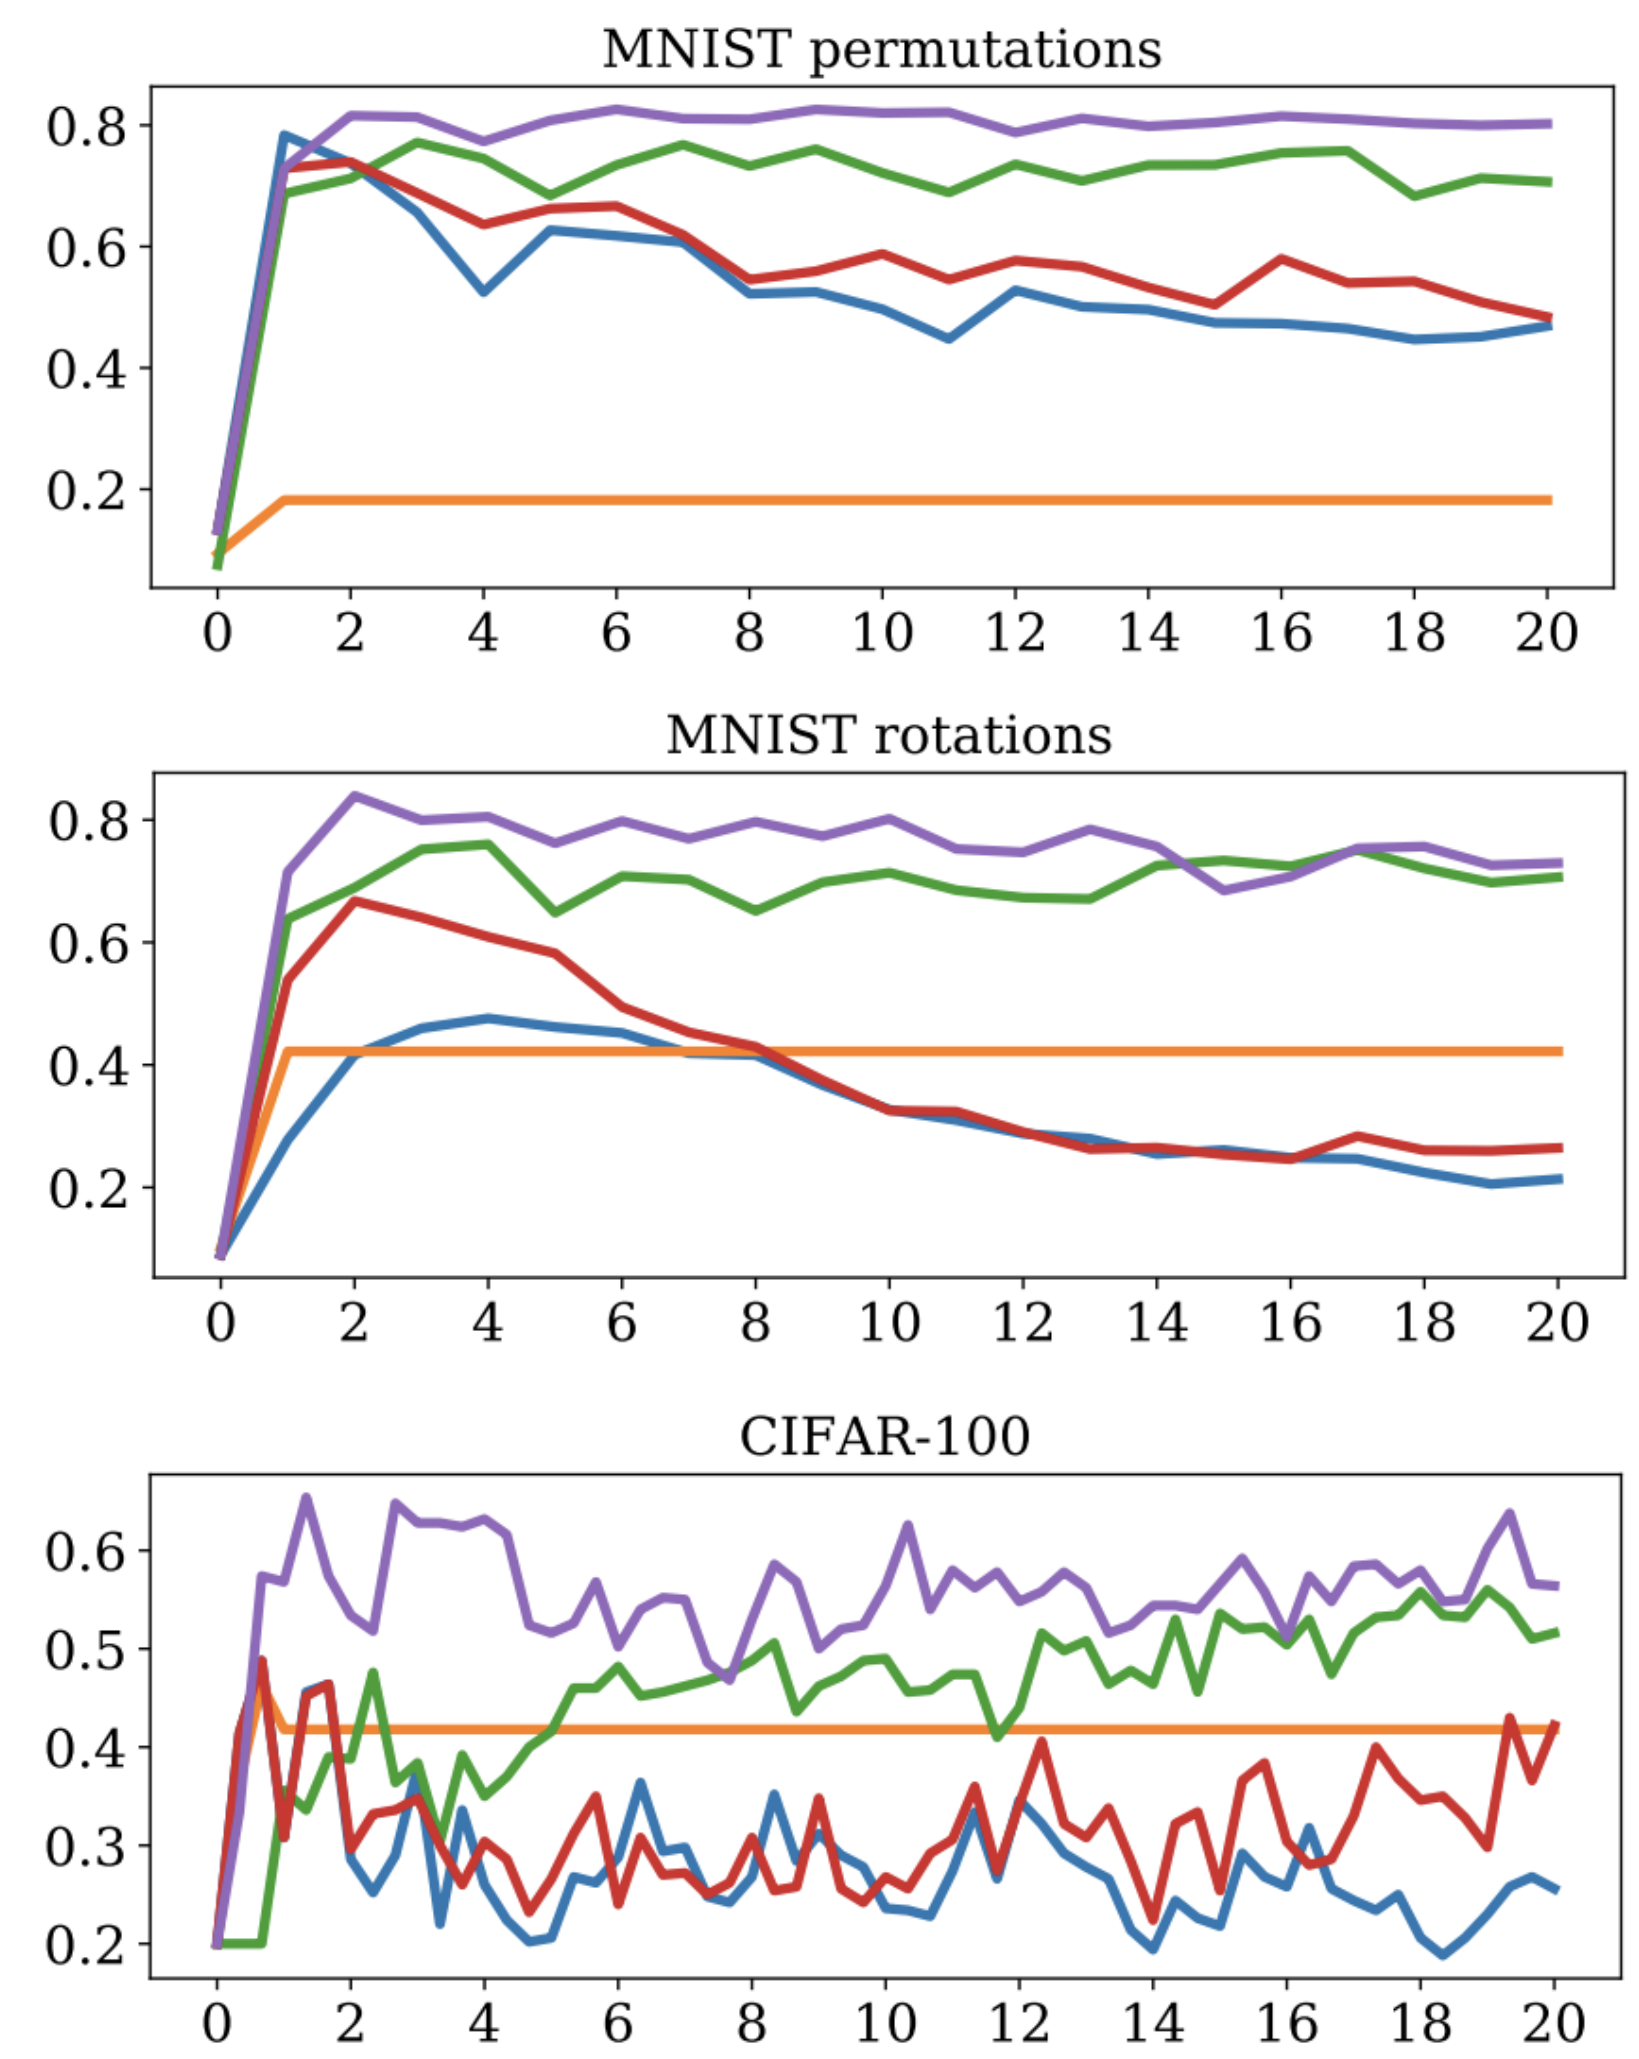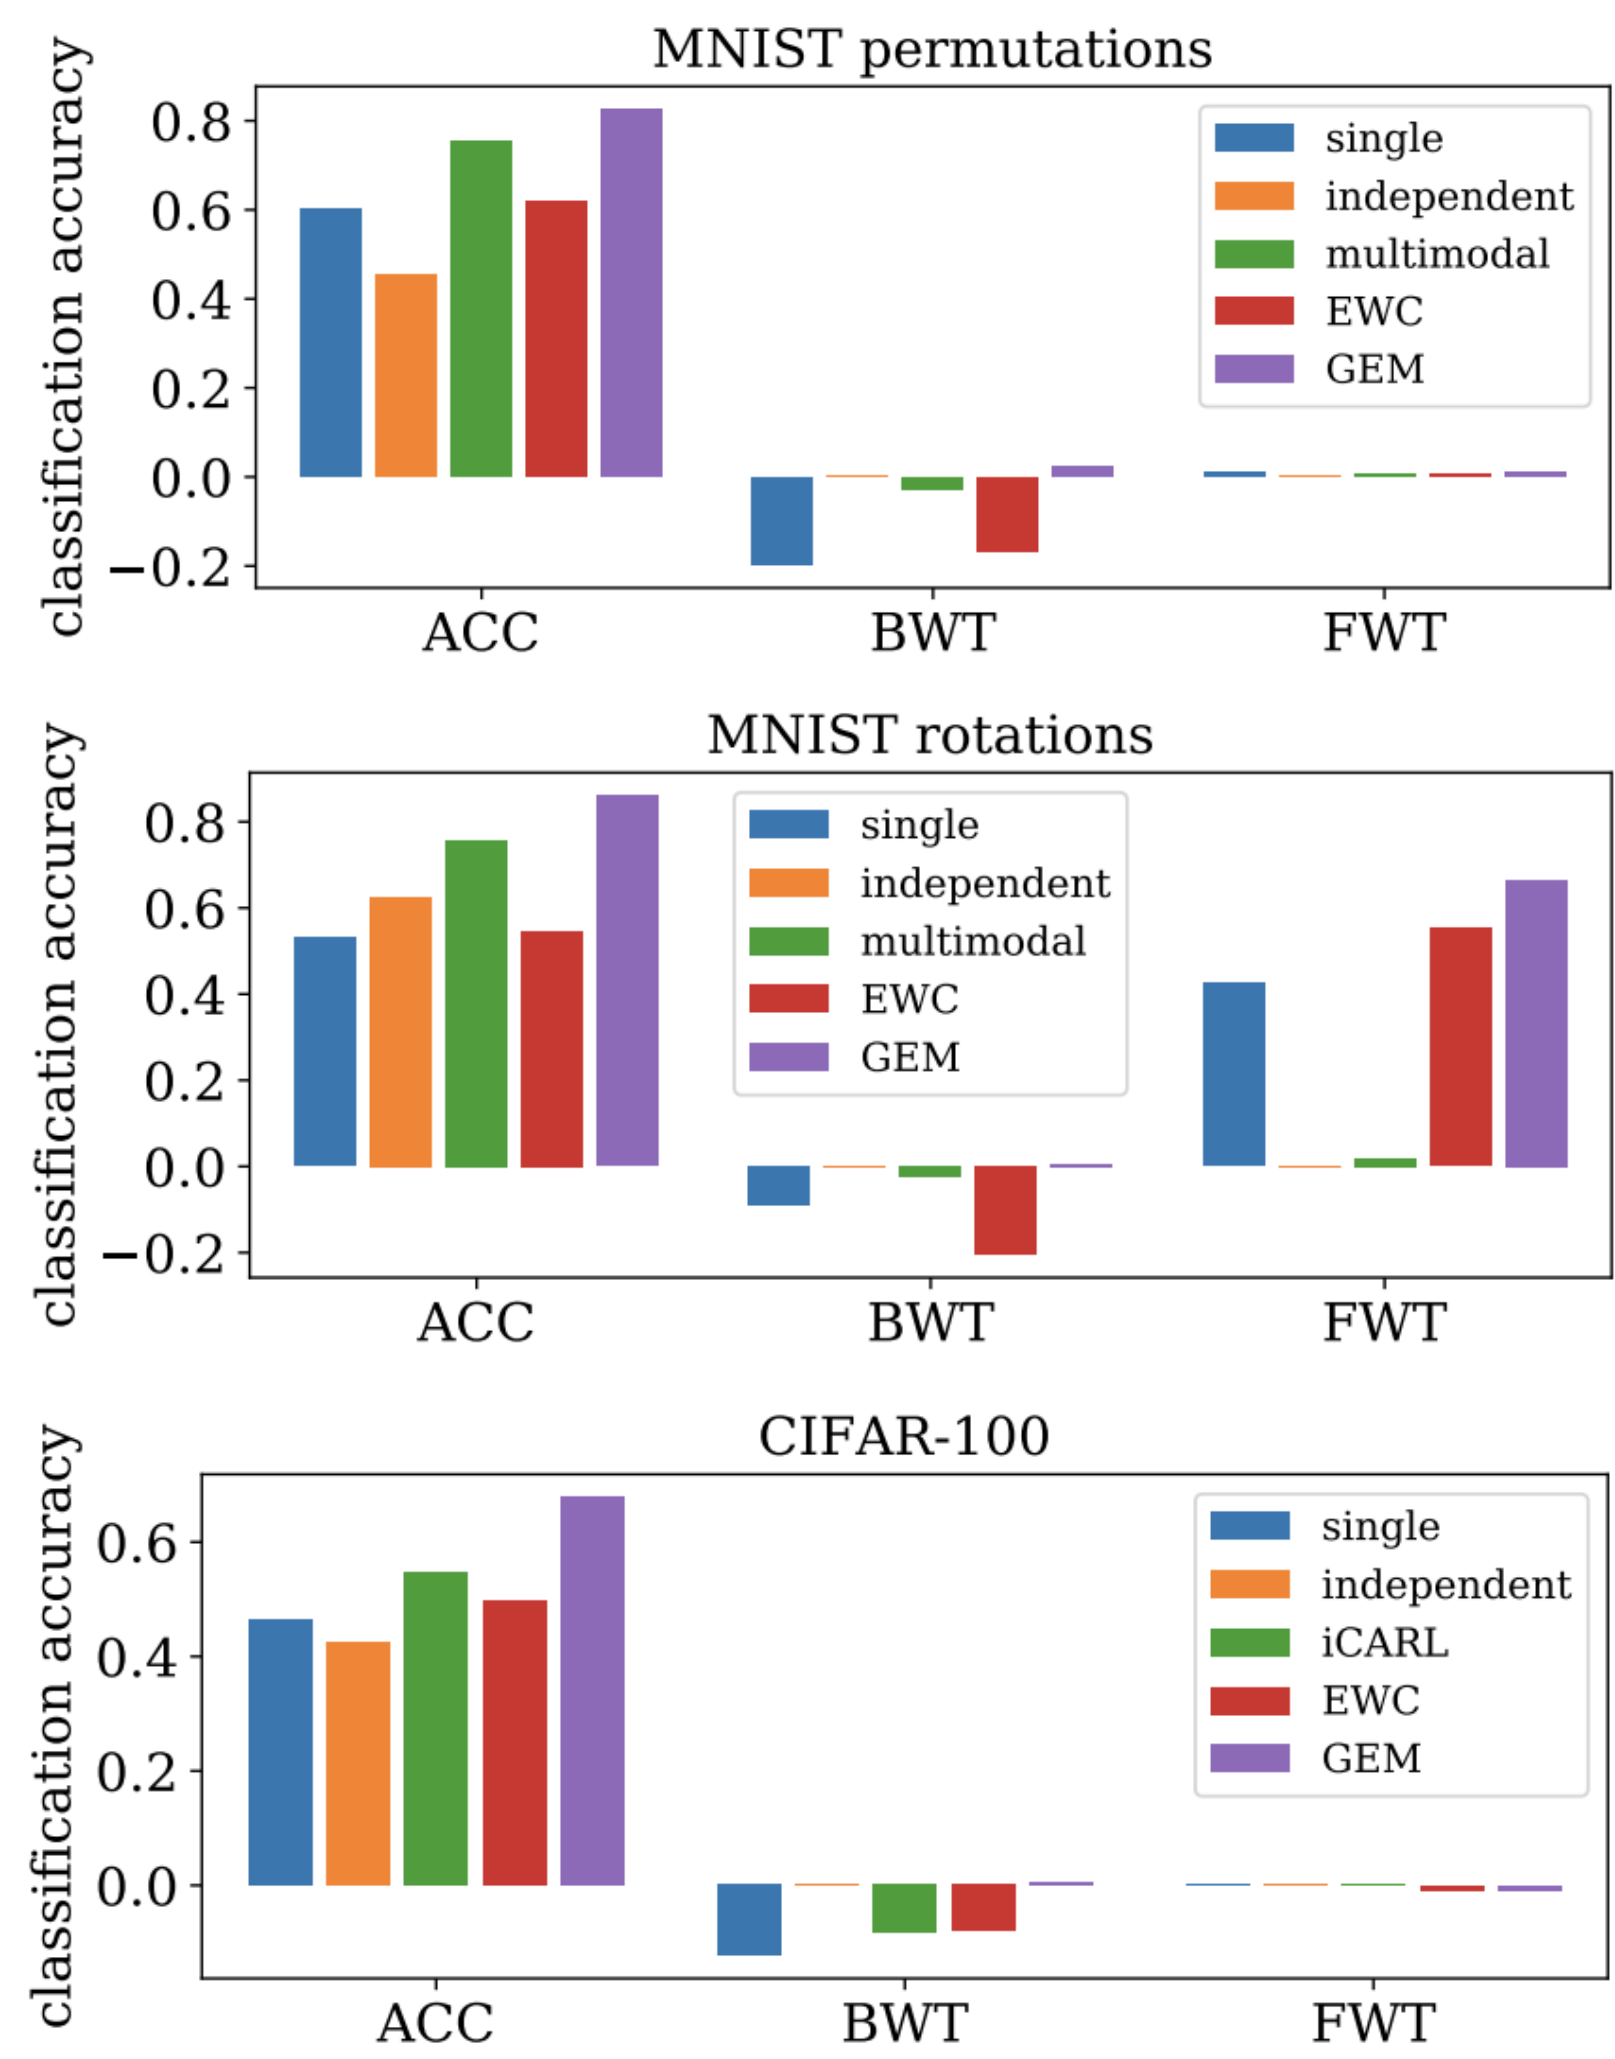Lopez-Paz & Ranzato. Gradient Episodic Memory for Continual Learning. NeurIPS '17

Experiments

Problems:
- MNIST permutations
- MNIST rotations
- CIFAR-100 (5 new classes/task)

BWT: backward transfer,
FWT: forward transfer

Total memory size:
5012 examples

If we take a step back…   do these experimental domains make sense?

Lopez-Paz & Ranzato. Gradient Episodic Memory for Continual Learning. NeurIPS '17

Can we meta-learn how to avoid negative backward transfer?

Javed & White. *Meta-Learning Representations for Continual Learning.* NeurIPS '19

Beaulieu et al. *Learning to Continually Learn.* '20

Plan for Today

The lifelong learning **problem statement**

Basic approaches to lifelong learning

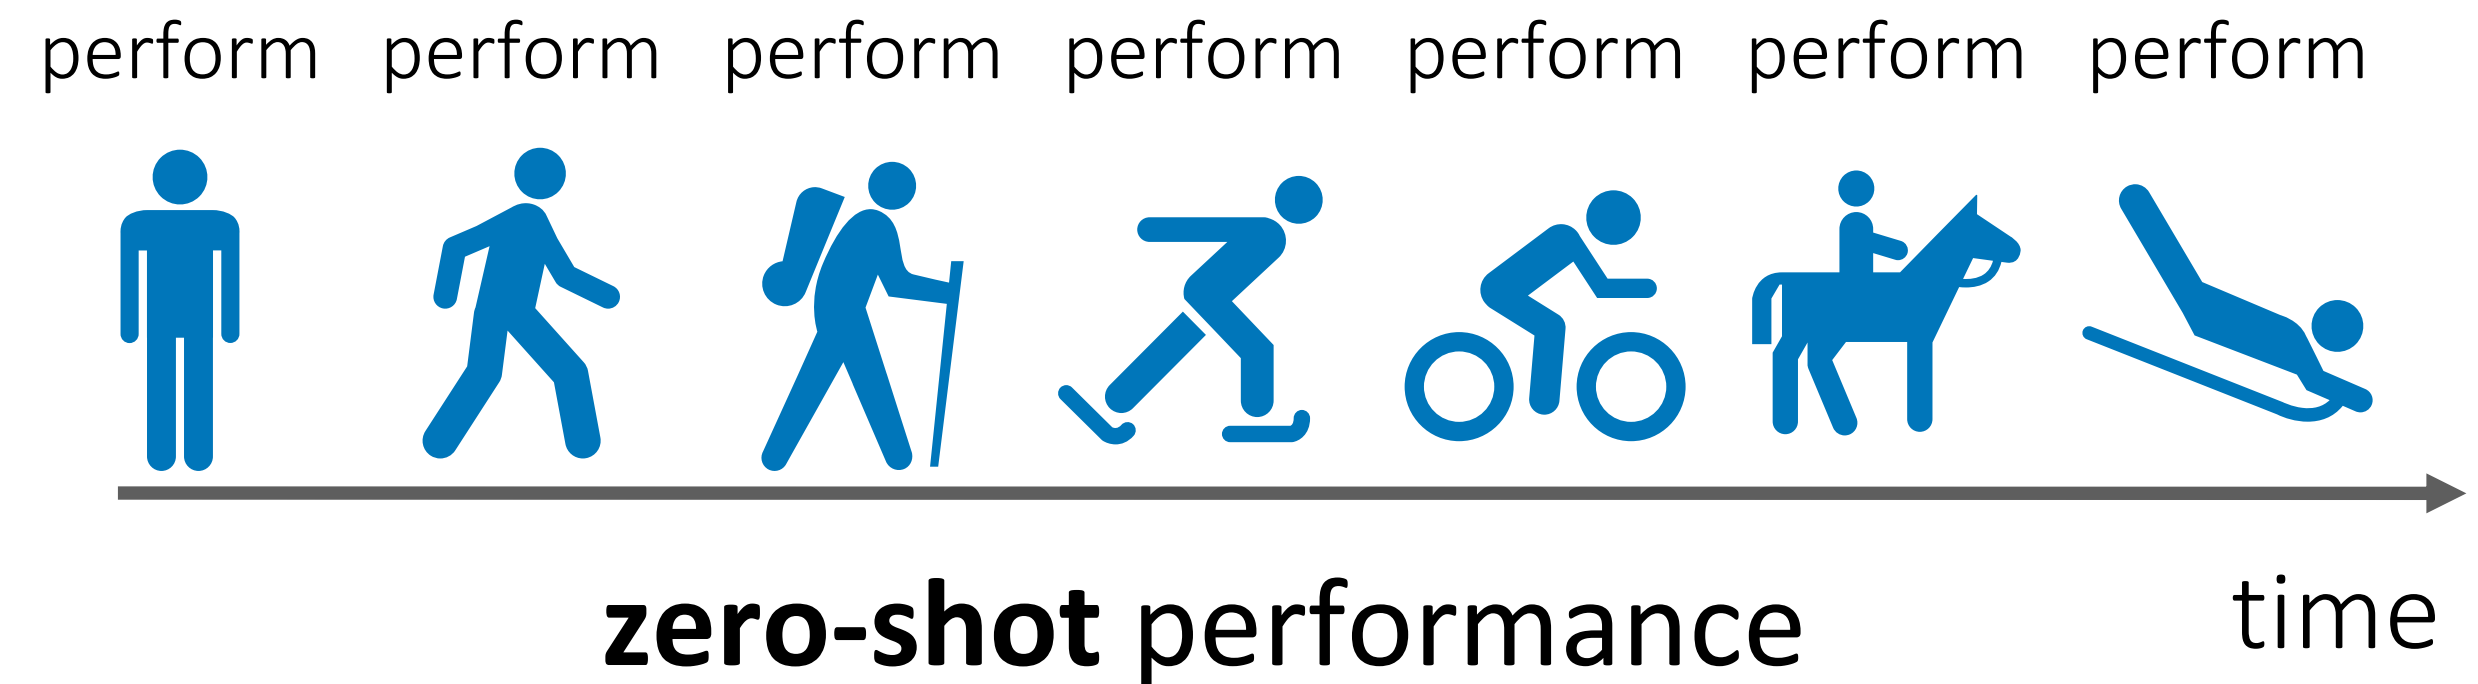Can we do **better** than the basics?

Revisiting the problem statement
from **the meta-learning perspective**

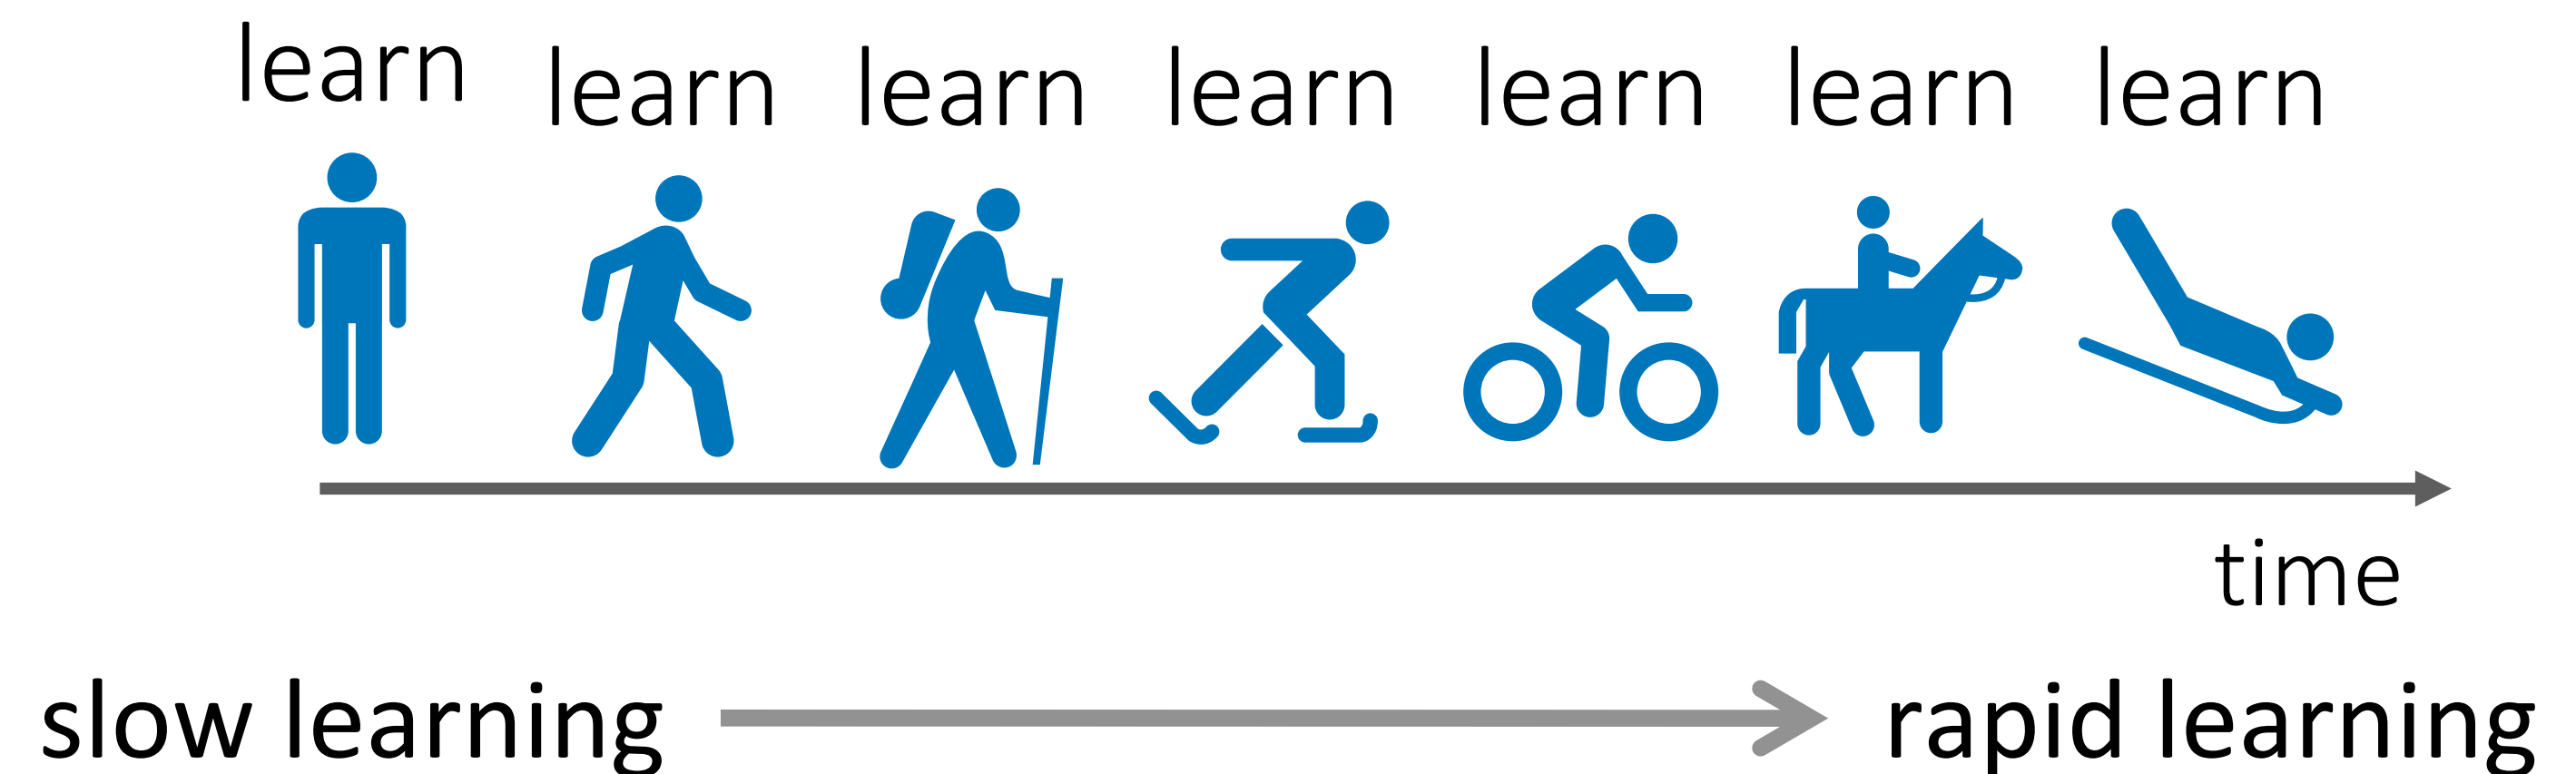# What might be wrong with the online learning formulation?

Online Learning
(Hannan '57, Zinkevich '03)

Perform sequence of tasks while minimizing static regret.

perform   perform   perform   perform   perform   perform   perform

zero-shot performance   time

More realistically:

learn   learn   learn   learn   learn   learn   learn
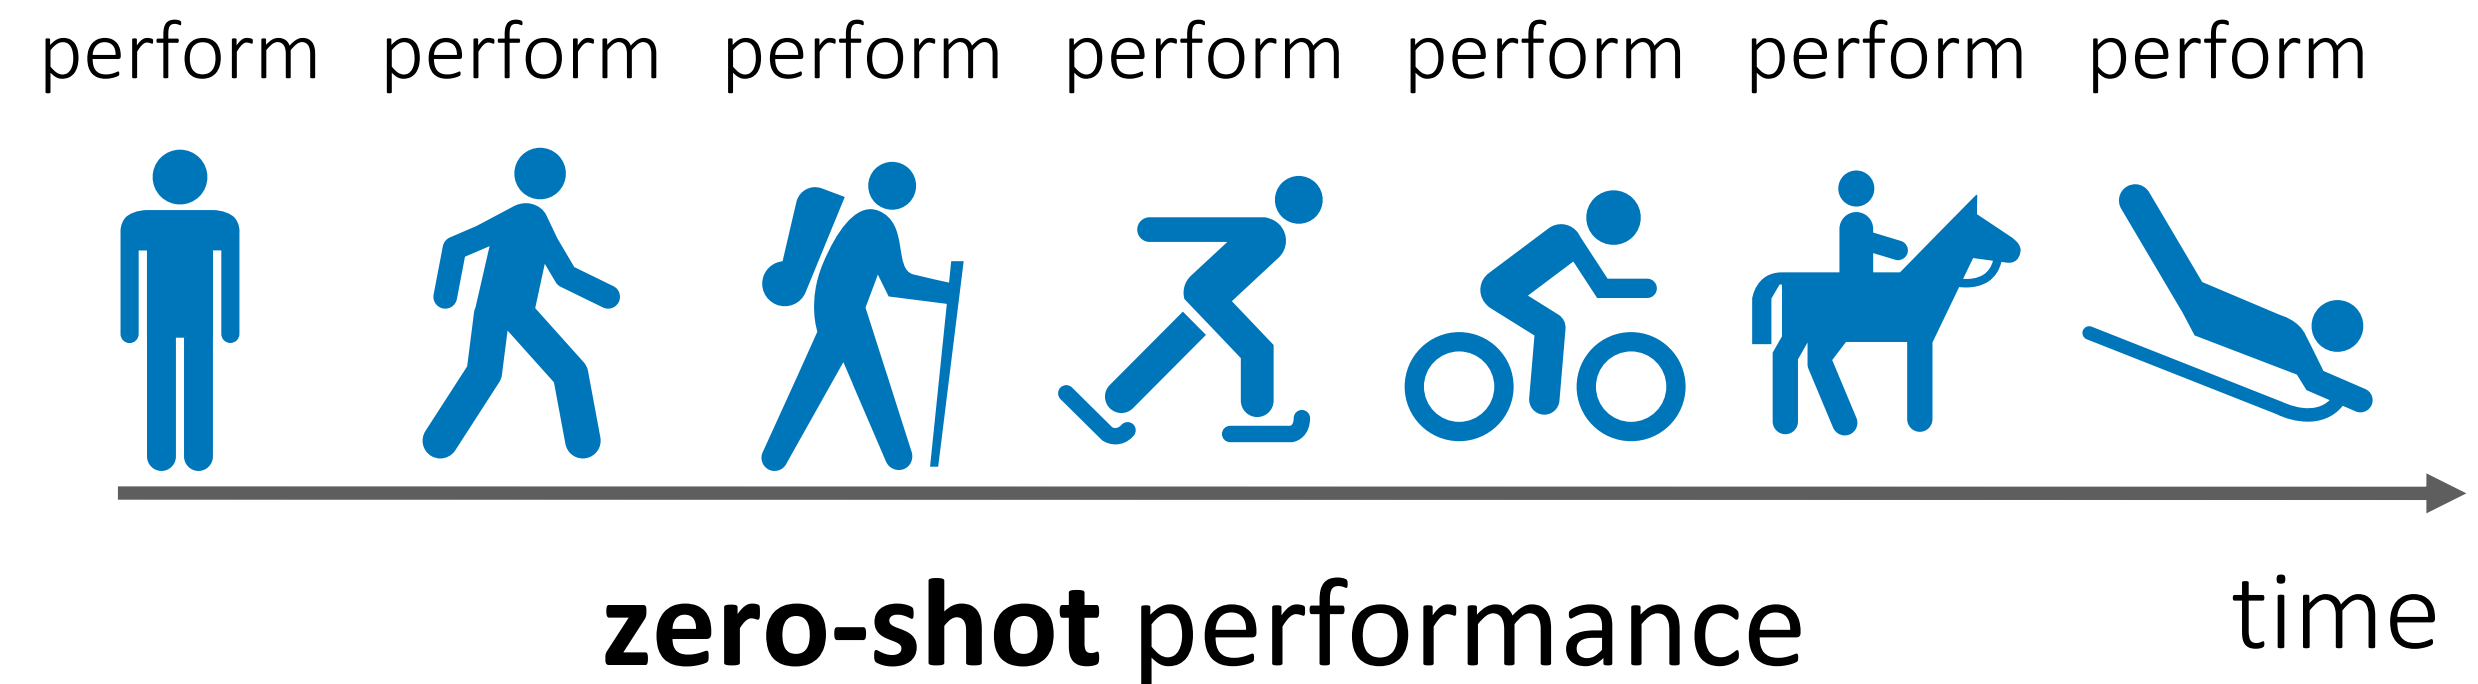
time

slow learning   →   rapid learning

What might be wrong with the online learning formulation?

Online Learning
(Hannan '57, Zinkevich '03)

Perform sequence of tasks while minimizing static regret.

perform perform perform perform perform perform perform

zero-shot performance   time

Online Meta-Learning

Efficiently learn a sequence of tasks from a non-stationary distribution.

learn learn learn learn learn learn learn

time

evaluate performance after seeing a small amount of data
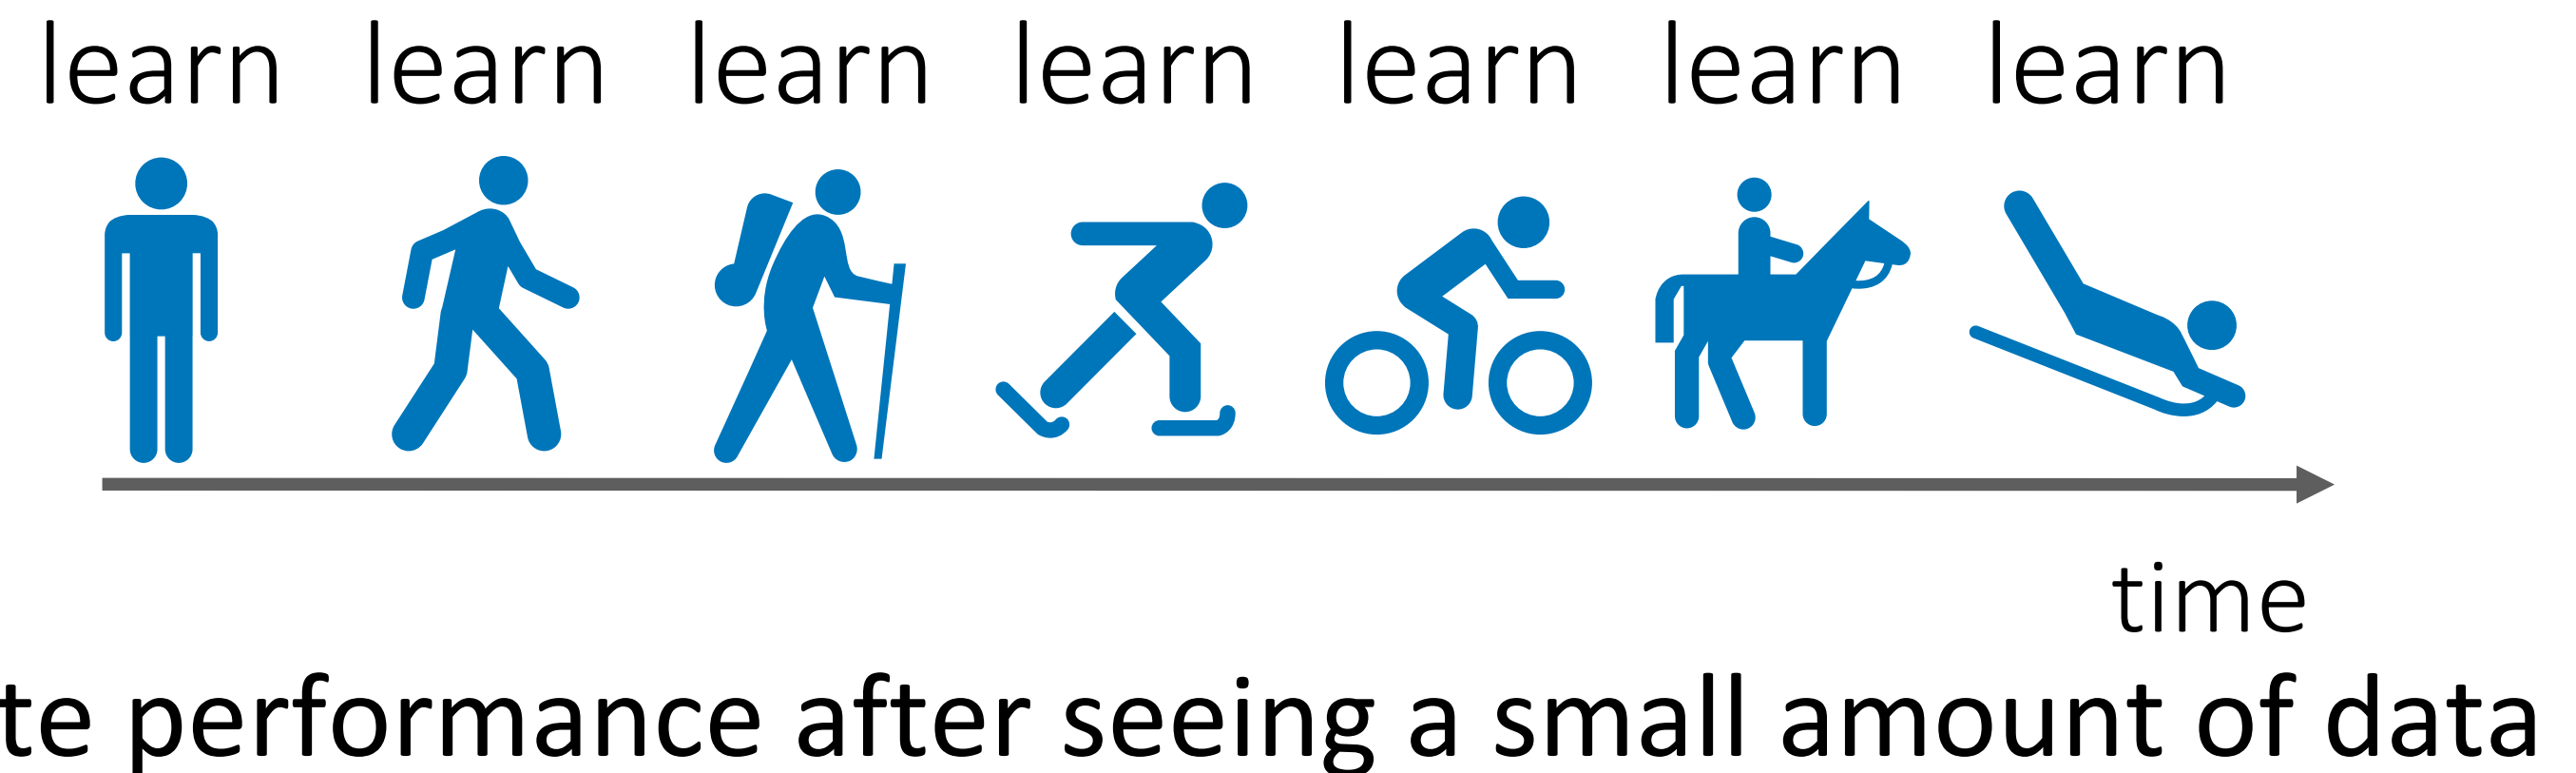
Primarily a difference in *evaluation*, rather than the *data stream*.

(Finn*, Rajeswaran*, Kakade, Levine ICML '18)

The Online Meta-Learning Setting

for task t = 1, ..., n

$\quad$ observe $\mathcal{D}_t^{tr}$

$\quad$ use update procedure $\Phi(\theta_t, \mathcal{D}_t^{tr})$ to produce parameters ϕ_t

$\quad$ observe x_t

$\quad$ predict $\hat{y}_t = f_{\phi_t}(x_t)$ $\qquad$ Standard online learning setting

$\quad$ observe label y_t

Loss of algorithm $\qquad$ Loss of best algorithm in hindsight

Goal: Learning algorithm with sub-linear $\quad \mathrm{Regret}_T := \sum_{t=1}^{T} \ell_t(\Phi_t(\theta_t)) - \min_{\theta \in \Theta} \sum_{t=1}^{T} \ell_t(\Phi_t(\theta))$

(Finn*, Rajeswaran*, Kakade, Levine ICML '18)

38

Can we apply meta-learning in lifelong learning settings?

Recall the **follow the leader** (FTL) algorithm:

Store all the data you've seen so far, and train on it.

Deploy model on current task.

Follow the *meta*-leader (FTML) algorithm:

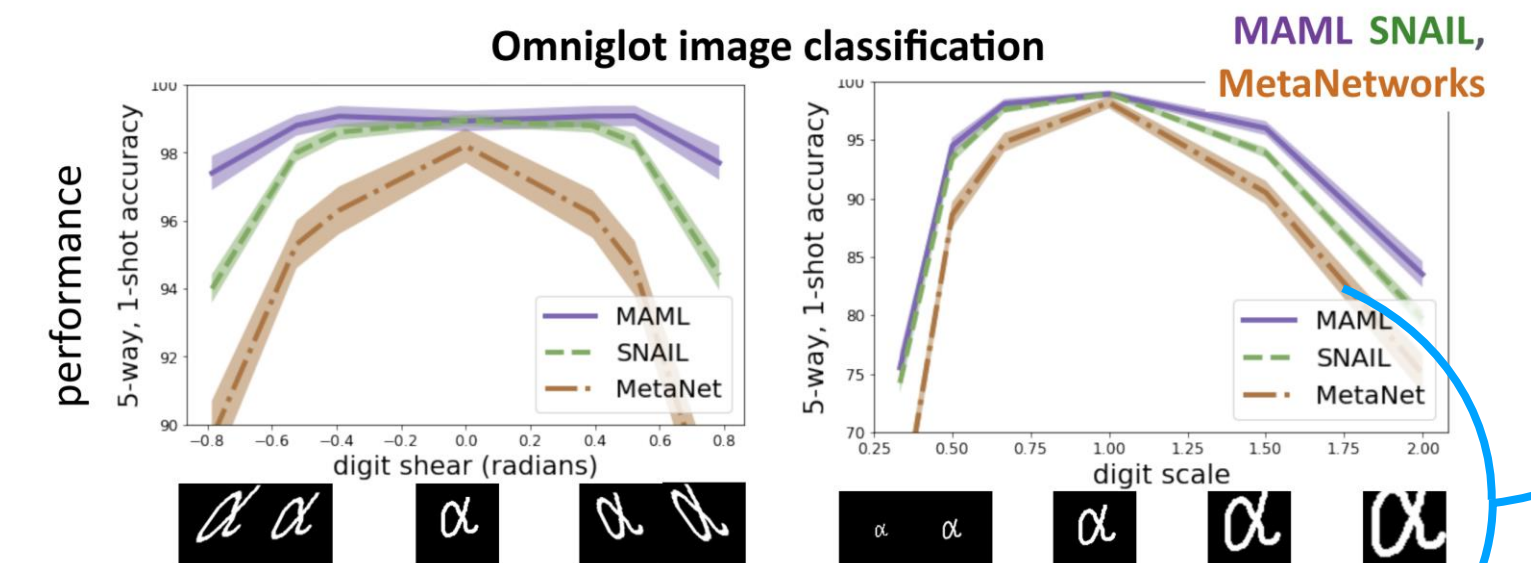Store all the data you've seen so far, and **meta**-train on it.

Run update procedure on the current task.

What meta-learning algorithms are well-suited for FTML?

What if $p_t(\mathcal{T})$ is non-stationary?

Optimization vs. Black-Box Adaptation

How well can learning procedures generalize to similar, but extrapolated tasks?

MAML SNAIL, MetaNetworks

Omniglot image classification

Does this structure come at a cost? Finn & Levine ICLR '18

Experiments

Experiment with **sequences of tasks**:
- Colored, rotated, scaled **MNIST**
- **3D object pose prediction**
- **CIFAR-100** classification
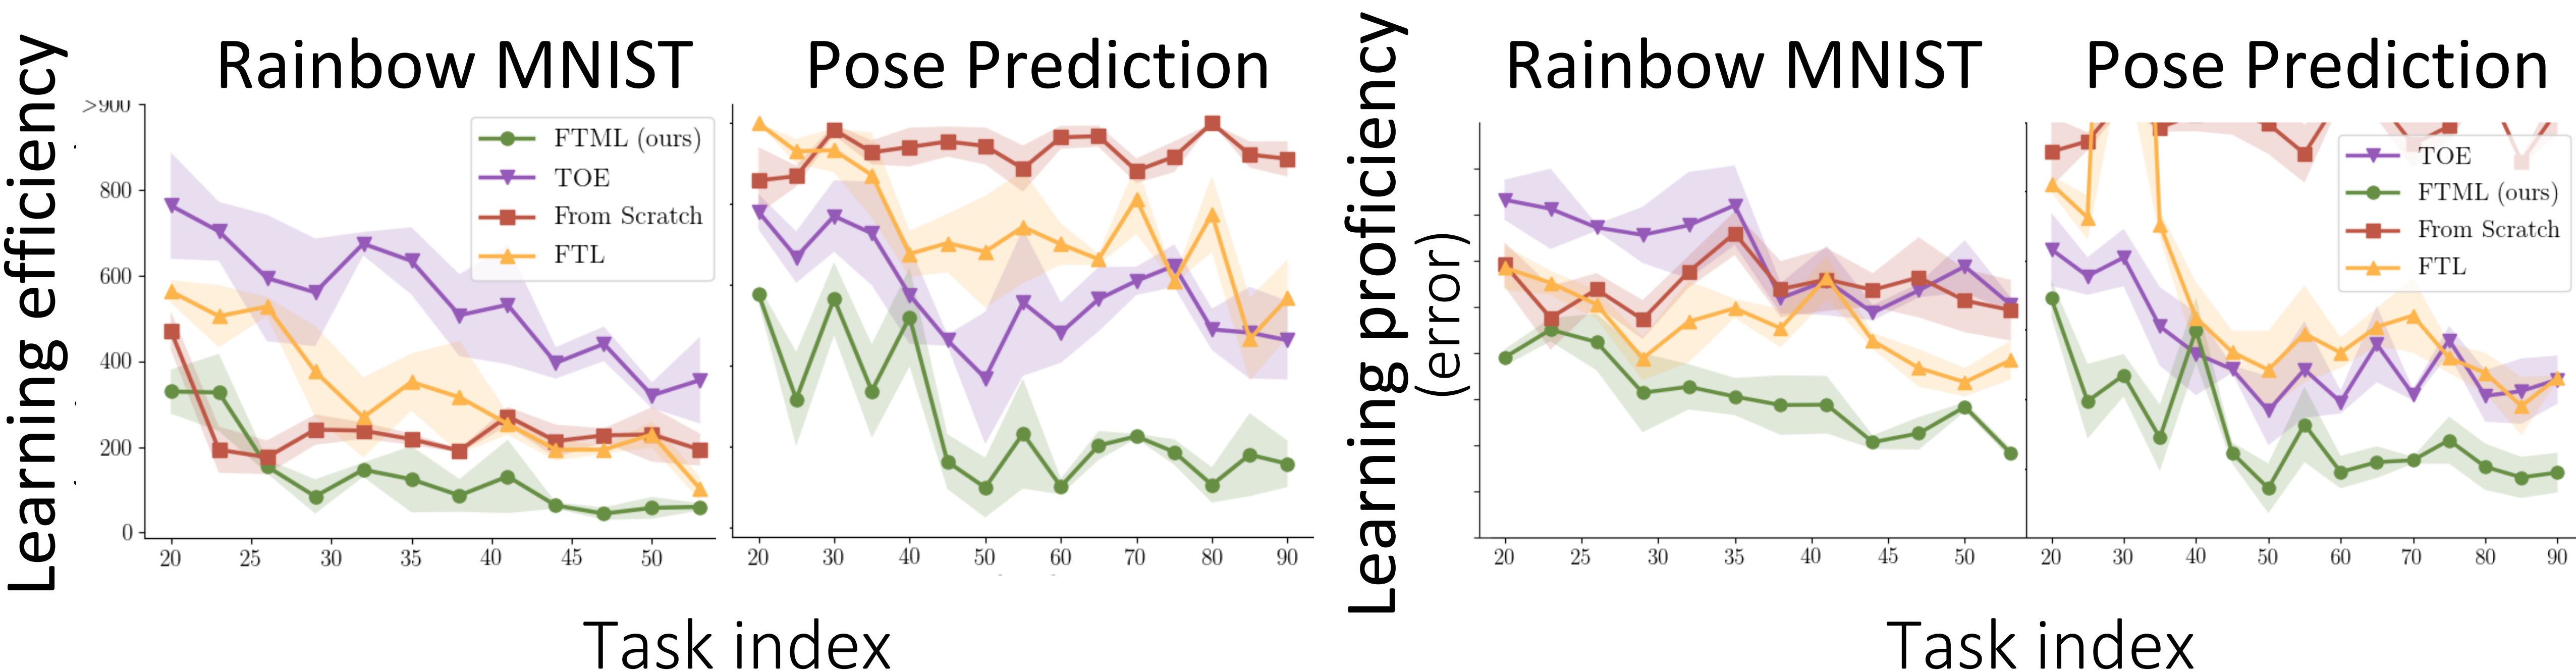
Example pose prediction tasks

plane

car

chair

Experiments

Comparisons:
- **TOE** (train on everything): train on all data so far
- **FTL** (follow the leader): train on all data so far, fine-tune on current task
- **From Scratch:** train from scratch on each task

**Follow The Meta-Leader
learns each new task faster** & **with greater proficiency**,
approaches **few-shot learning** regime

Takeaways

Many flavors of lifelong learning, all under the same name.

Defining the problem statement is often the hardest part

Meta-learning can be viewed as a slice of the lifelong learning problem.

A very open area of research.